Visualization of CNN

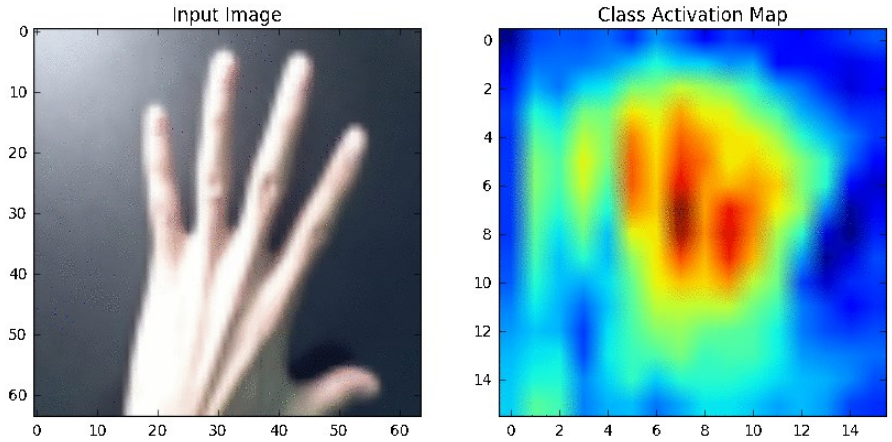

What's going on inside CNN?

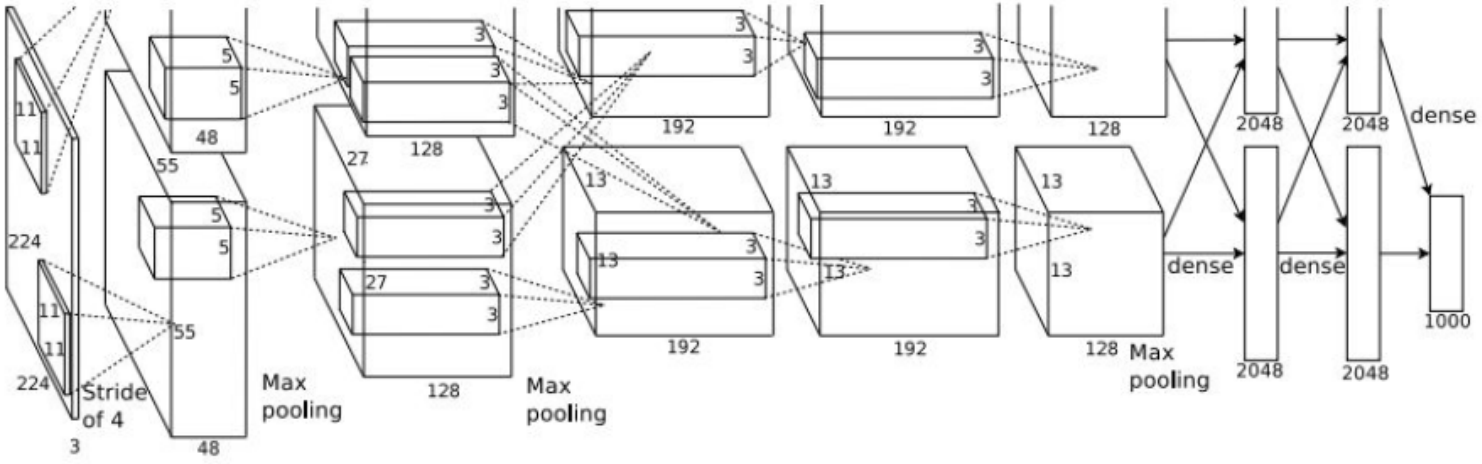

Class Scores:
1000 numbers

Input Image:
3 x 224 x 224

What are the intermediate features looking for?

Slide Credit : Stanford CS231n

Visualize Patches that Maximally Activate Neurons

Figure 4: Top regions for six pool$_5$ units. Receptive fields and activation values are drawn in white. Some units are aligned to concepts, such as people (row 1) or text (4). Other units capture texture and material properties, such as dot arrays (2) and specular reflections (6).

Rich feature hierarchies for accurate object detection and semantic segmentation
[Girshick, Donahue, Darrell, Malik]

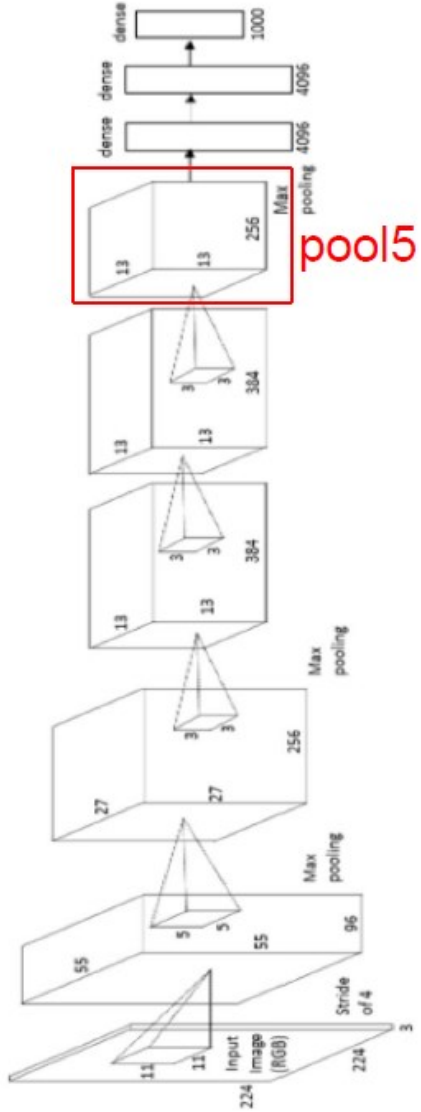

Slide Credit : Stanford CS231n

Visualize Filters

- Only interpretable on the first layer

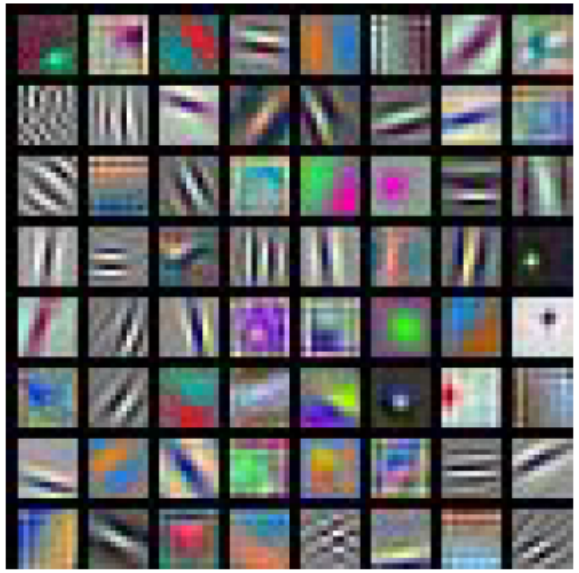

AlexNet:
64 x 3 x 11 x 11

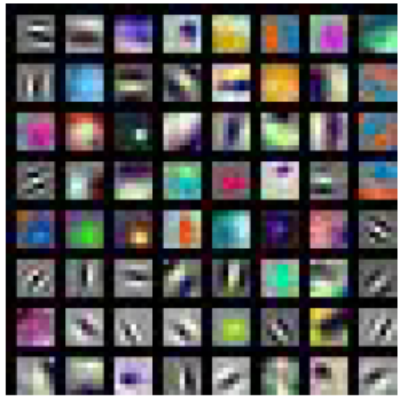

ResNet-18:
64 x 3 x 7 x 7

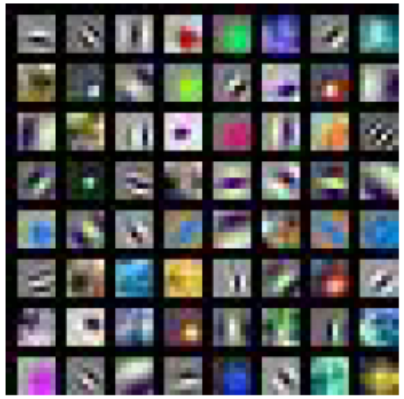

ResNet-101:
64 x 3 x 7 x 7

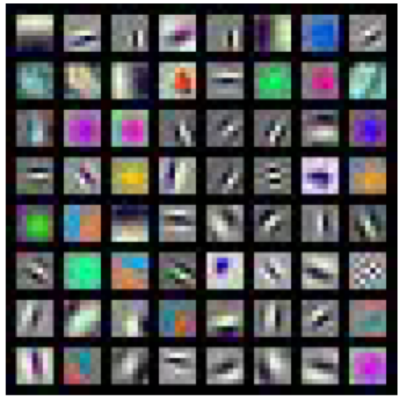

DenseNet-121:
64 x 3 x 7 x 7

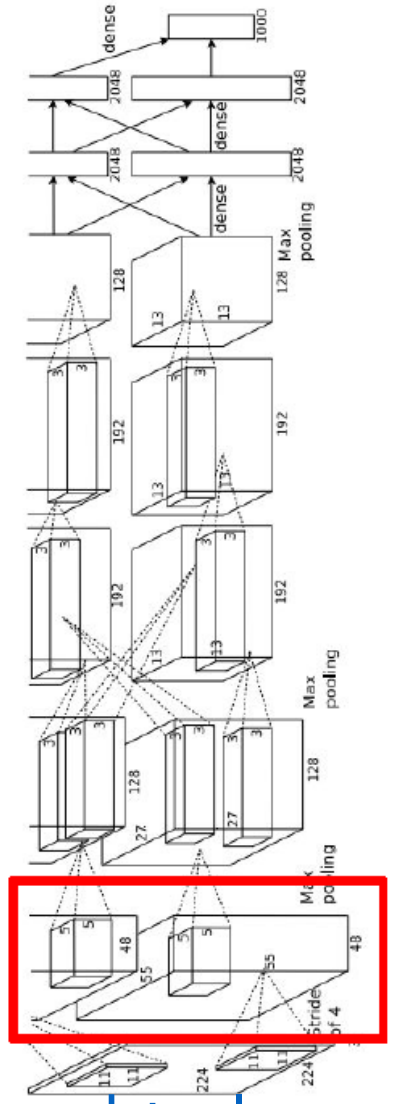

- http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html

Slide Credit : Stanford CS231n

Visualizing Activations

- https://www.youtube.com/watch?v=AgkfIQ4IGaM

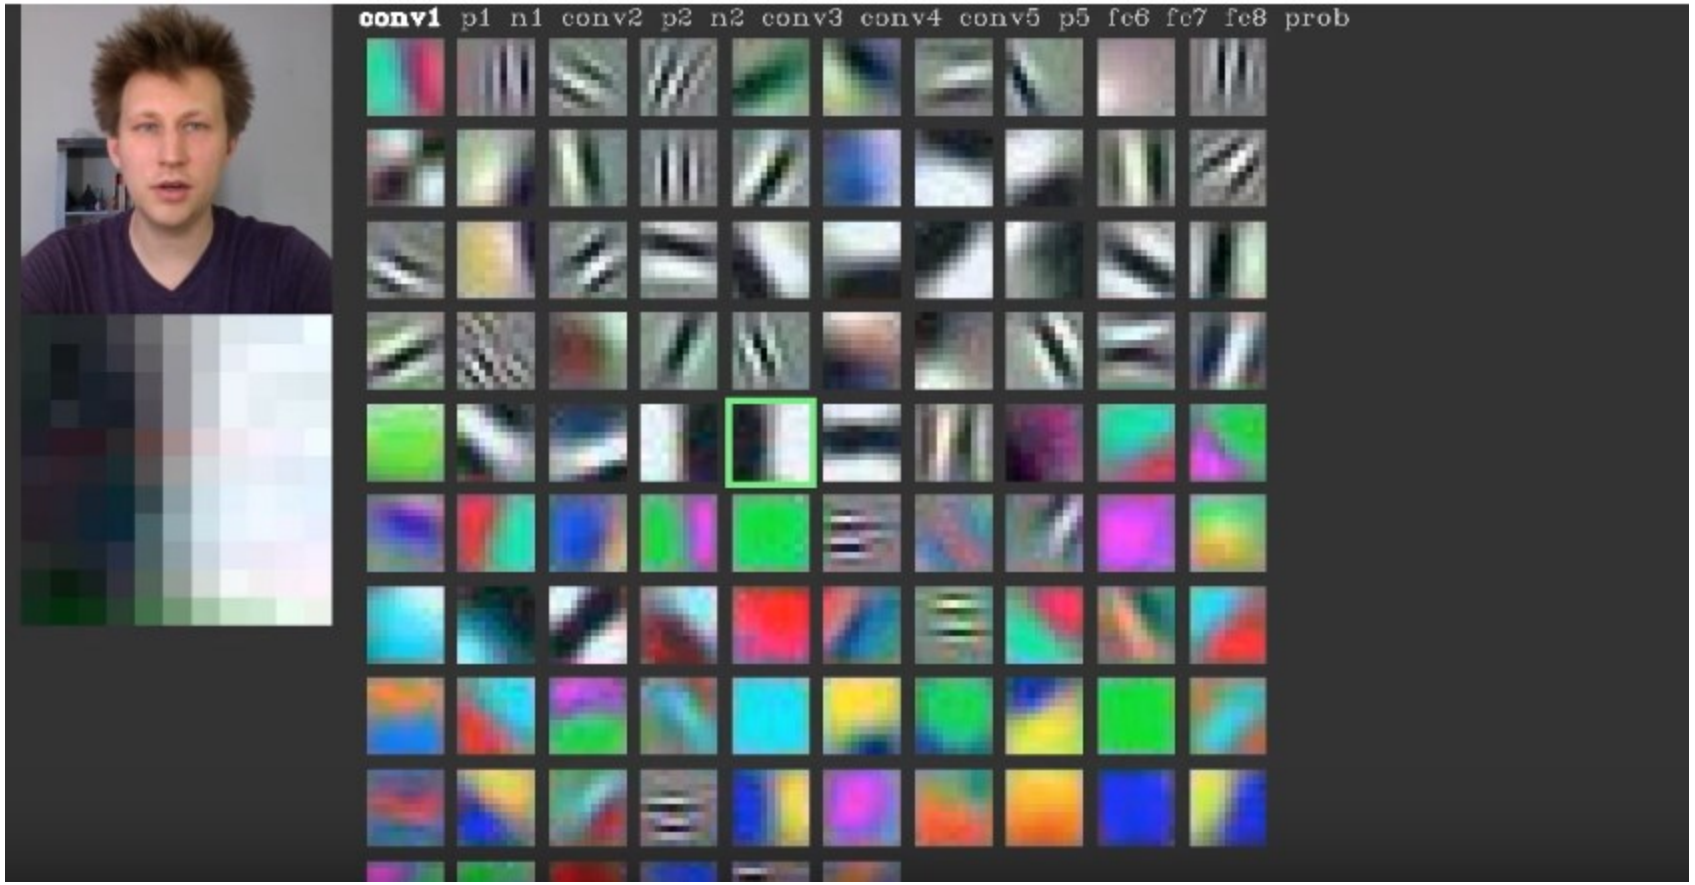

Visualizing the Representation

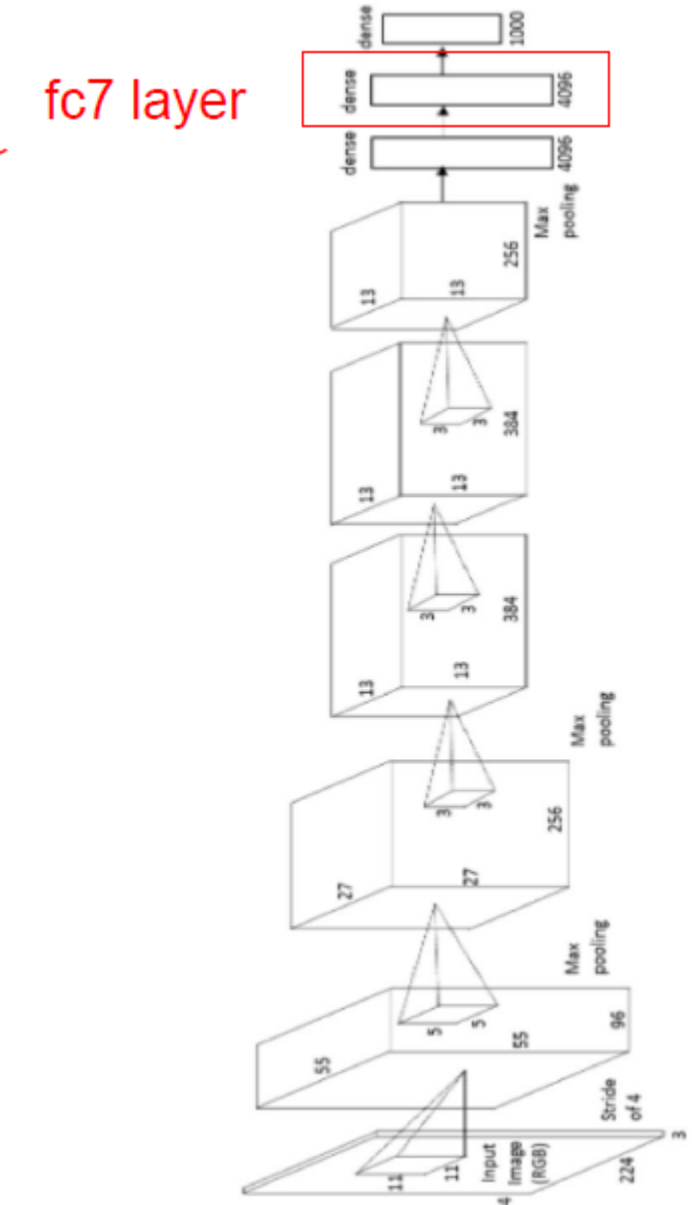

4096-dimensional “code” for an image
(layer immediately before the classifier)

can collect the code for many images

Slide Credit : Stanford CS231n

Last Layer : Nearest Neighbors

4096-dim vector

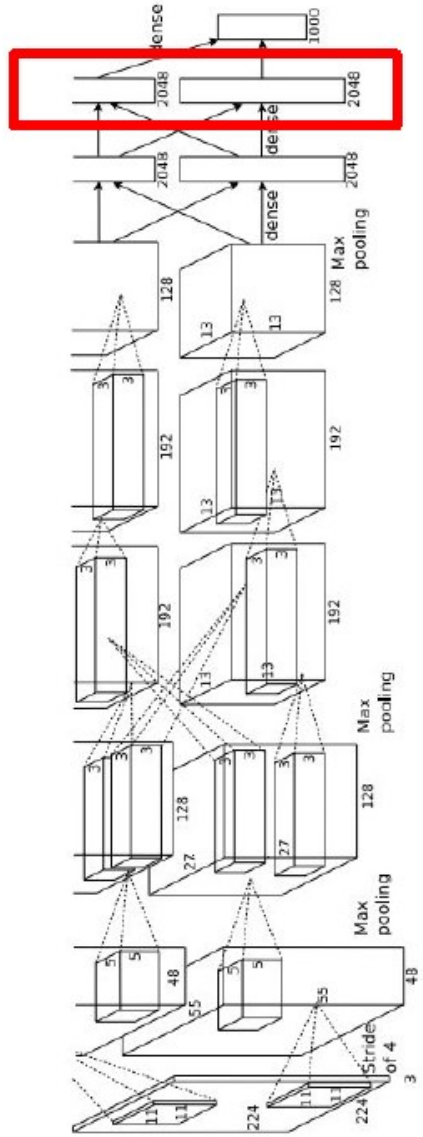

Test image

L2 Nearest neighbors in feature space

Recall: Nearest neighbors in pixel space

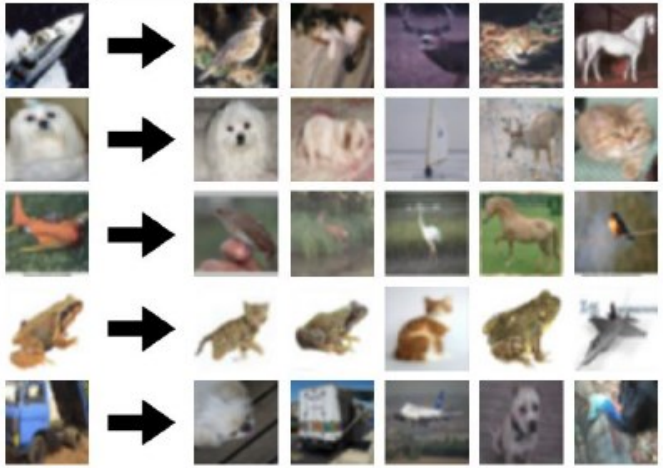

Last Layer: Dimensionality Reduction

- Visualize the "space" of FC7 feature vectors by reducing dimensionality of vectors from 4096 to 2 dimensions
- Simple algorithm: Principle Component Analysis(PCA)
- More complex: t-SNE

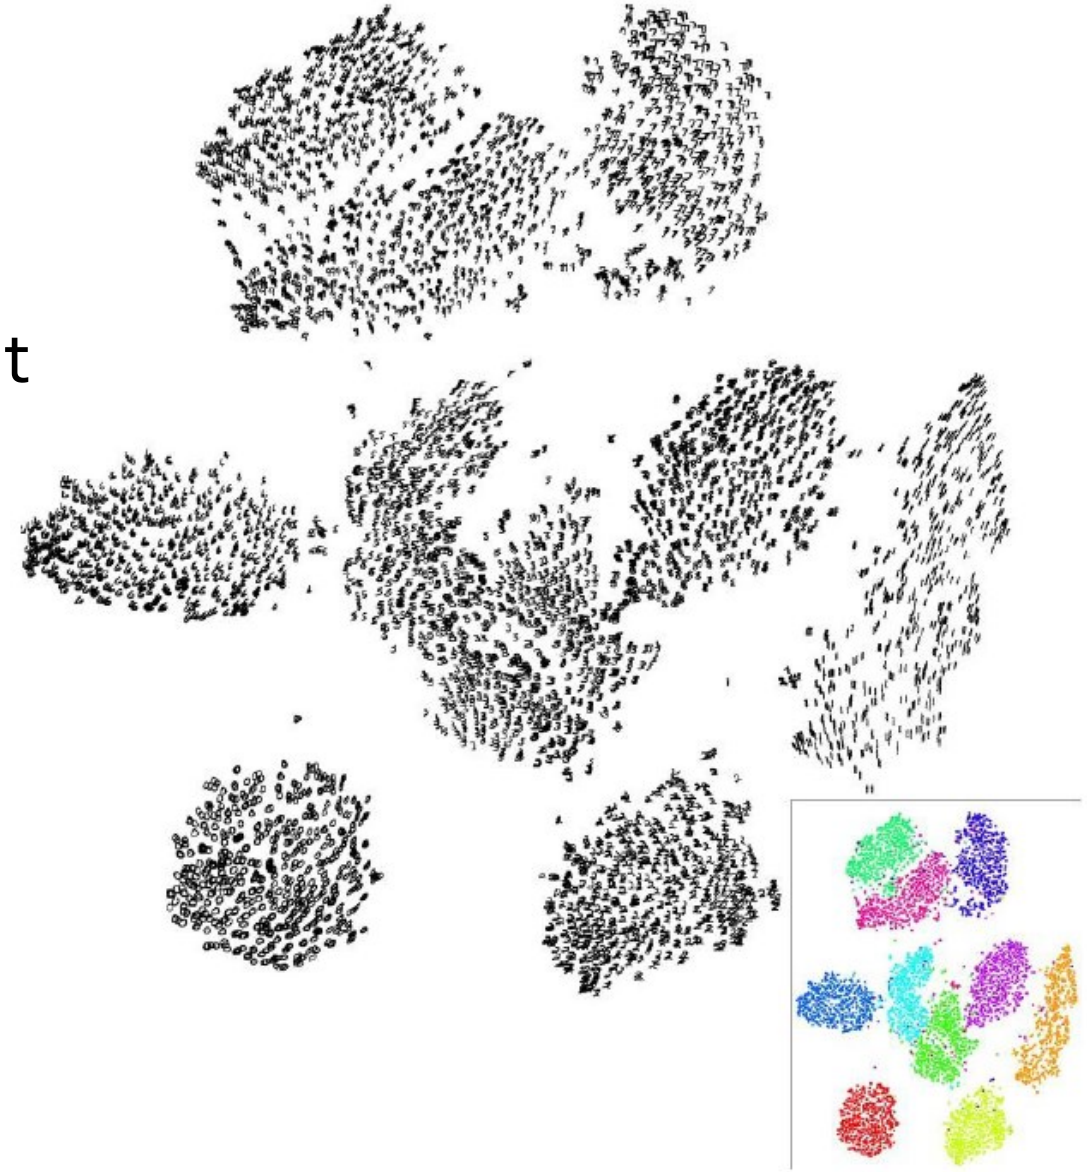

Last Layer: Dimensionality Reduction

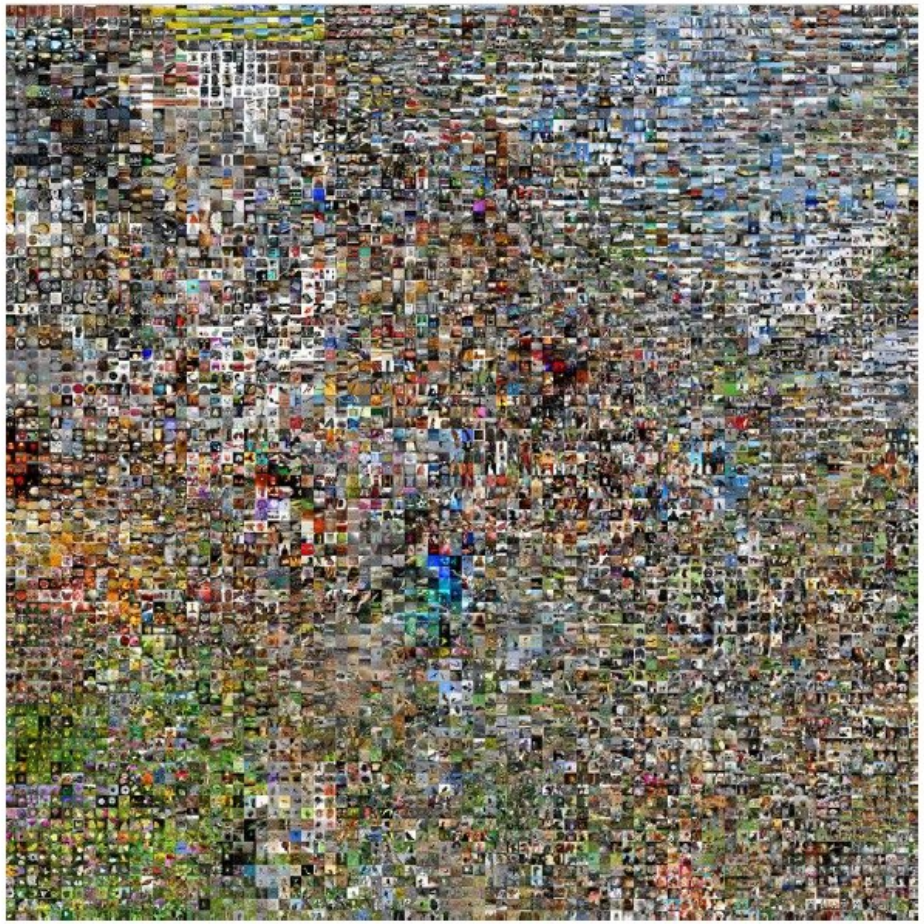

- http://cs.stanford.edu/people/karpathy/cnnembed/

Occlusion Experiments

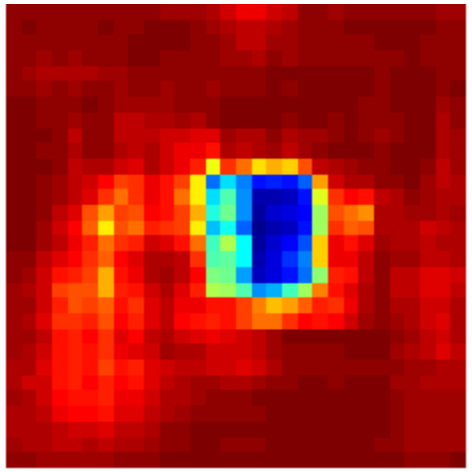

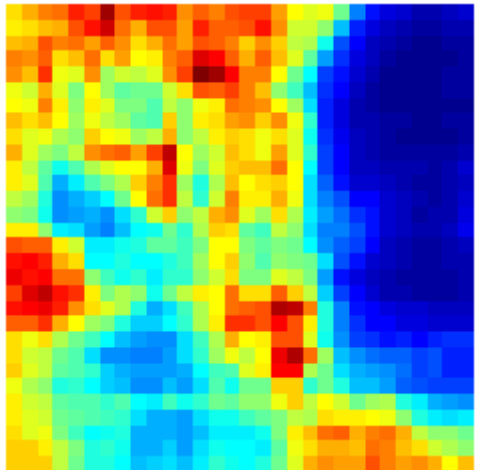

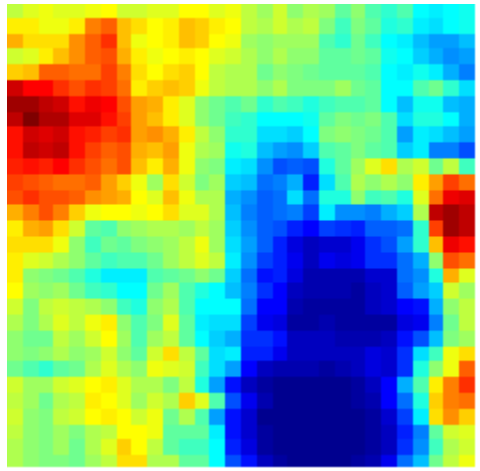

Weakly Supervised Learning

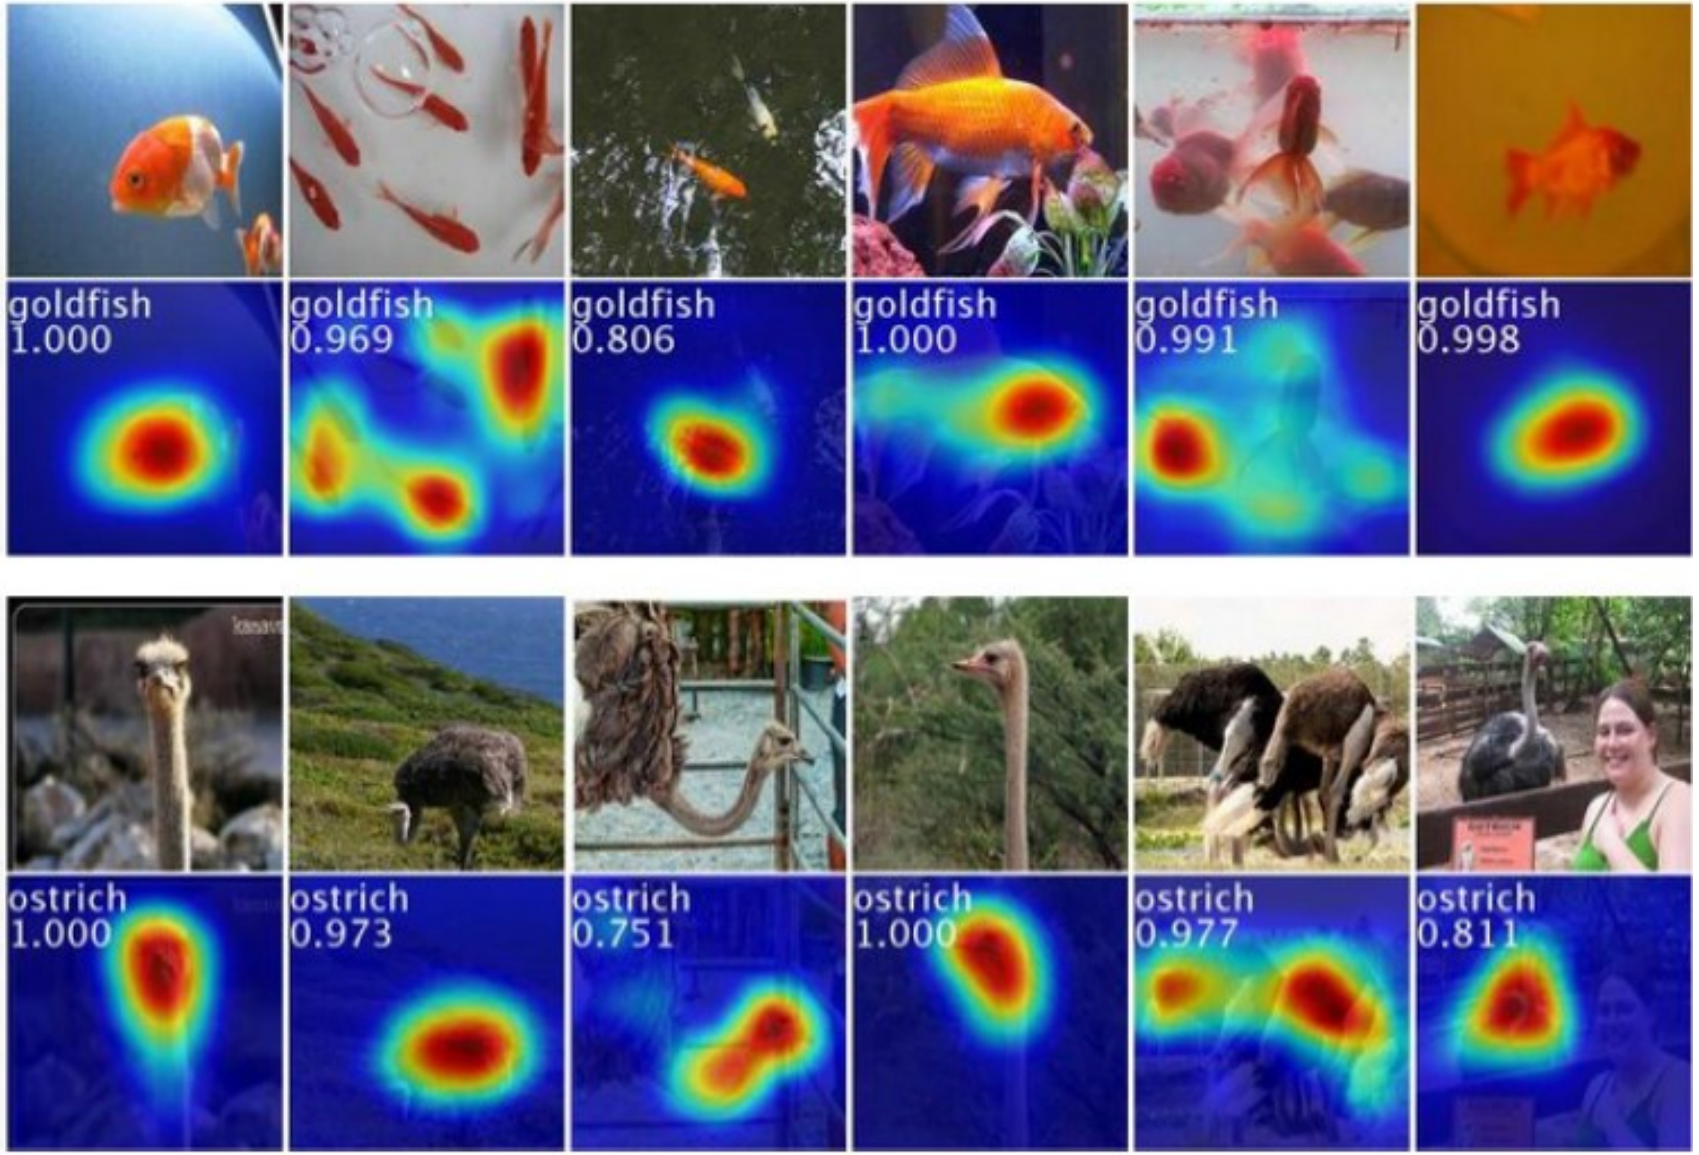

Class activation map (CAM)

- **Identify important image regions** by projecting back the weights of output layer to convolutional feature maps.
- CAMs can be generated for each class in single image.
- Regions for each categories are different in given image.
  - palace, dome, church …

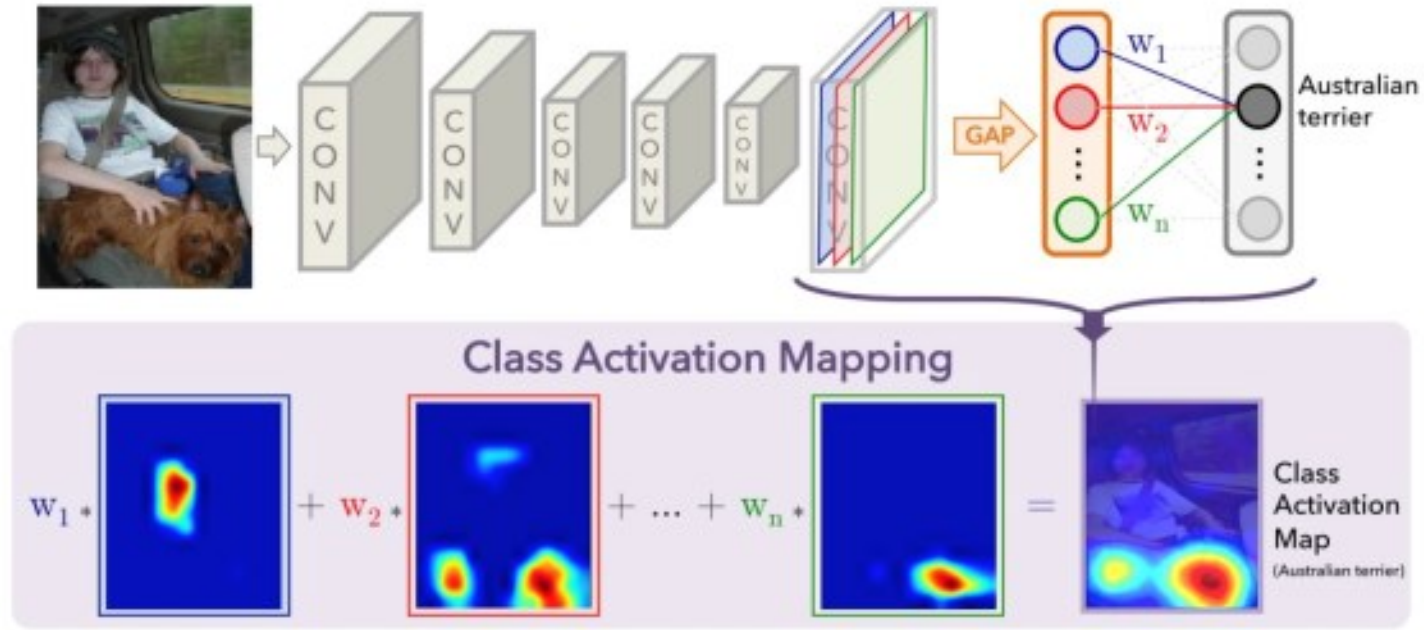

Figure 2. Class Activation Mapping: the predicted class score is mapped back to the previous convolutional layer to generate the class activation maps (CAMs). The CAM highlights the class-specific discriminative regions.

Results

- CAM on top 5 predictions on an image
- CAM for one object class in images

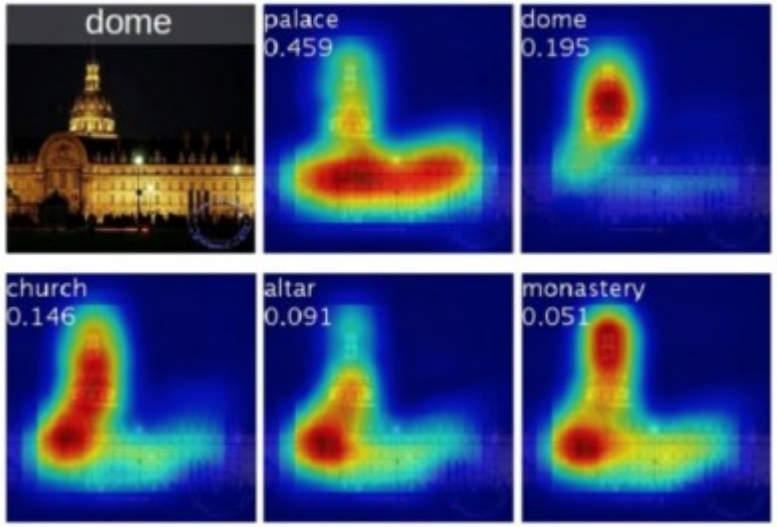

Figure 4. Examples of the CAMs generated from the top 5 predicted categories for the given image with ground-truth as dome. The predicted class and its score are shown above each class activation map. We observe that the highlighted regions vary across predicted classes e.g., *dome* activates the upper round part while *palace* activates the lower flat part of the compound.

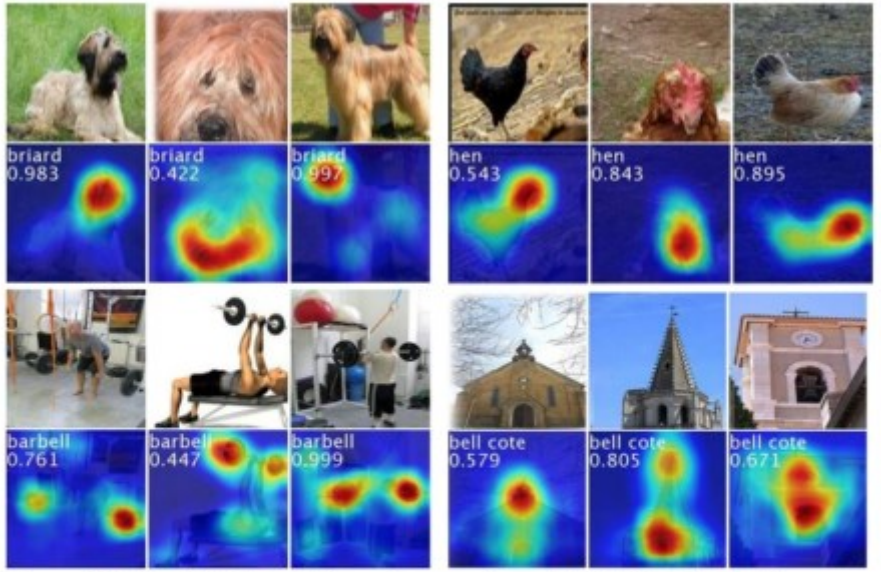

Figure 3. The CAMs of four classes from ILSVRC [20]. The maps highlight the discriminative image regions used for image classification e.g., the head of the animal for *briard* and *hen*, the plates in *barbell*, and the bell in *bell cote*.

Weakness of CAM (Weakly Supervised Localicztion)

- Focusing on discriminative features

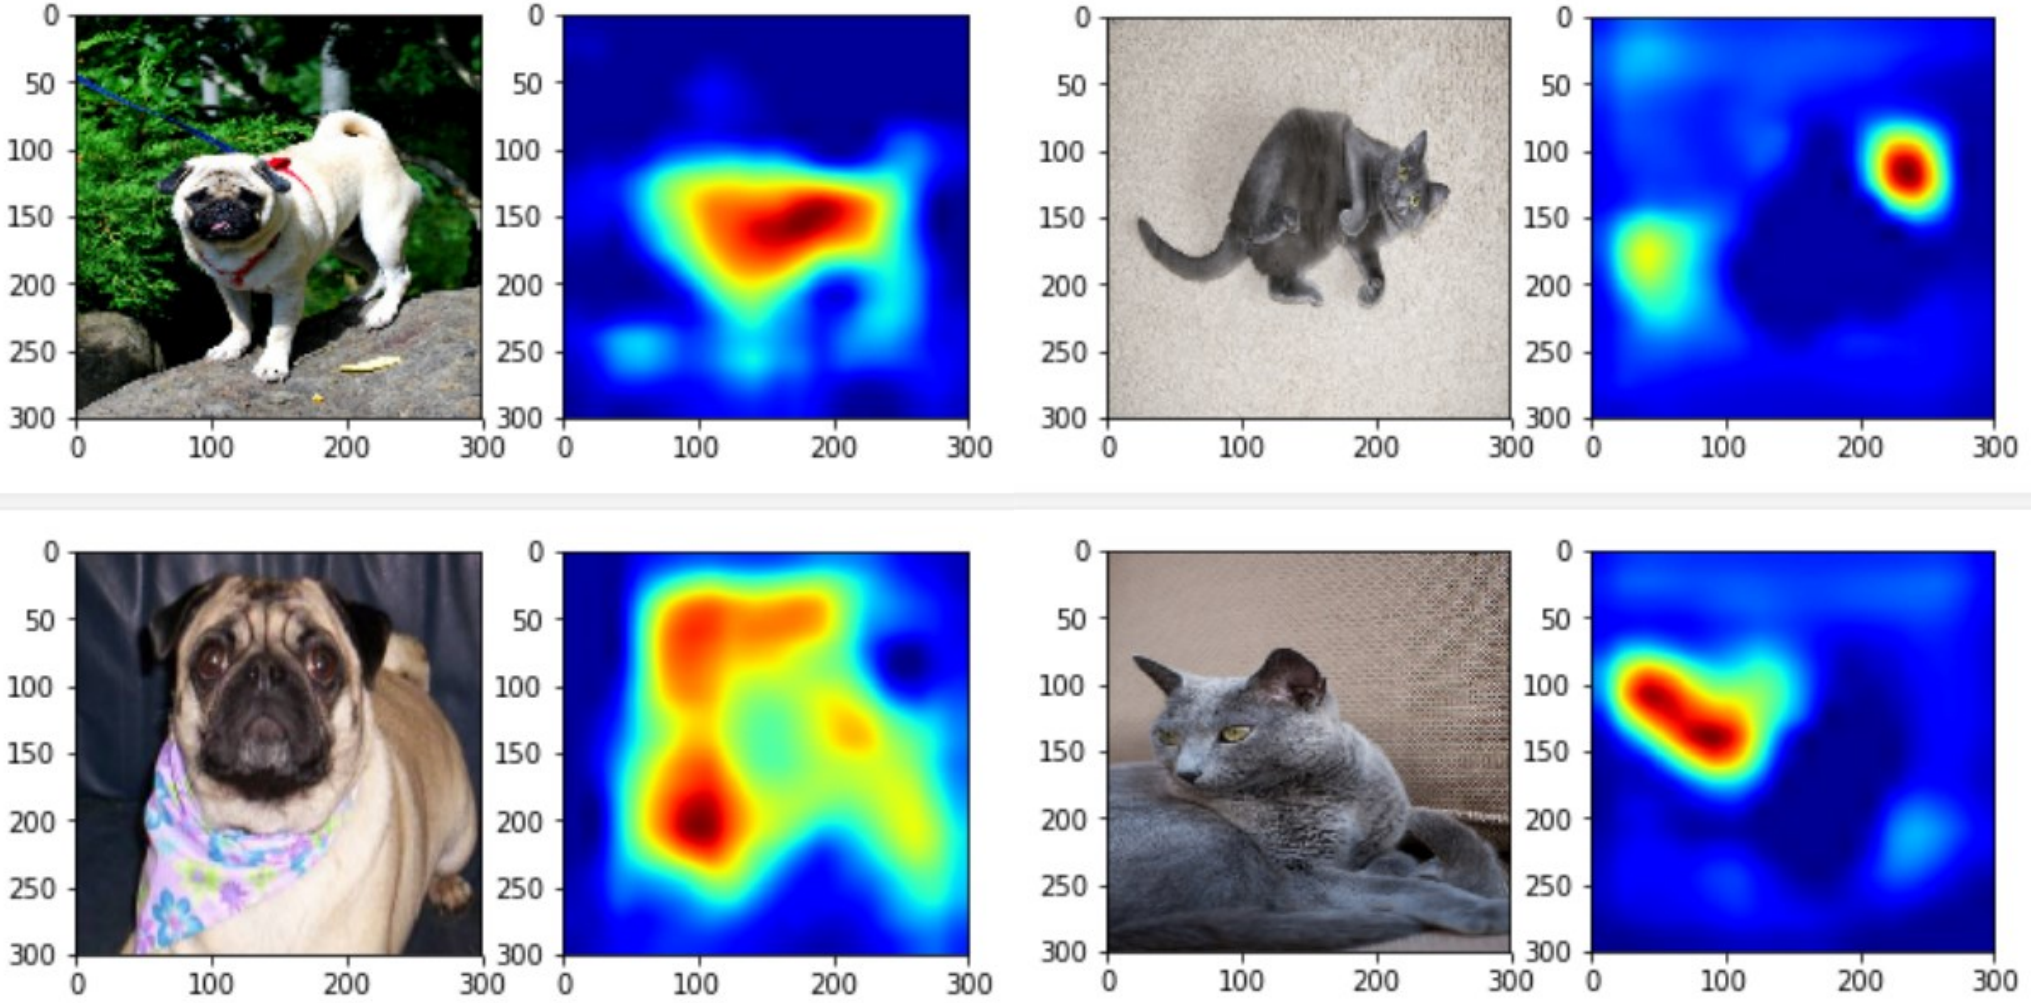

Weakness of CAM (Weakly Supervised Localicztion)

- Focusing on discriminative features